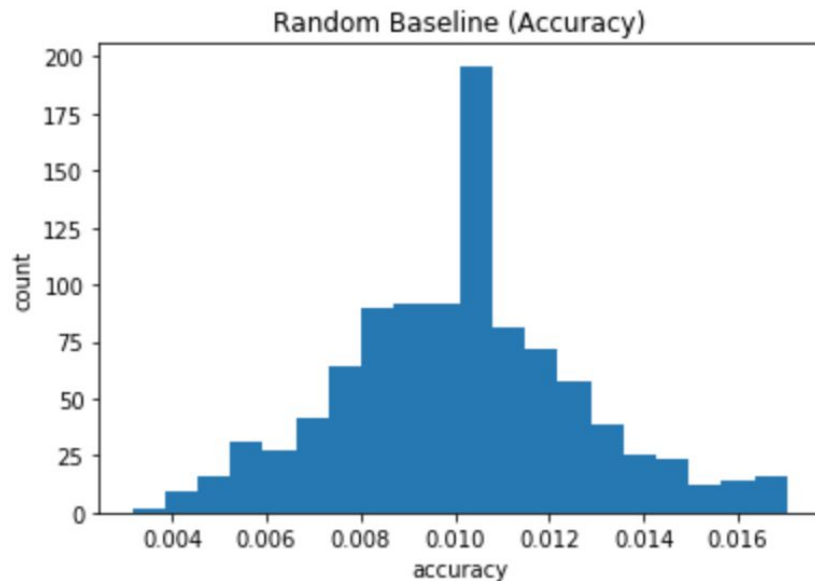


Landmark recognition

Introduction

We'll train a landmark recognition model using the Google Landmark Dataset containing 4,132,914 images belonging to 203,094 classes. We scale the data down to ~81,300 images belonging to 100 classes. And since the image is our only feature we'll skip the feature selection part.



Random Baseline(100 classes)



Average Accuracy : 1.01%

Three-Layer Feedforward Neural Networks

Accuracy on validation : $0.1 < \%$





Common schemes

- INPUT -> [CONV -> RELU -> POOL]*N -> [FC -> RELU]*M -> FC
- INPUT -> [CONV -> RELU -> CONV -> RELU -> POOL]*N -> [FC -> RELU]*M -> FC
- If you working with images, just use a winning architecture.

Ref: https://github.com/ekapolc/pattern_2019



CNN with Maxpool

Layer (type)	Output Shape	Param #
conv2d_10 (Conv2D)	(None, 297, 297, 100)	2800
max_pooling2d_8 (MaxPooling2)	(None, 148, 148, 100)	0
conv2d_11 (Conv2D)	(None, 146, 146, 100)	90100
max_pooling2d_9 (MaxPooling2)	(None, 73, 73, 100)	0
conv2d_12 (Conv2D)	(None, 71, 71, 100)	90100
max_pooling2d_10 (MaxPooling)	(None, 35, 35, 100)	0
flatten_4 (Flatten)	(None, 122500)	0
dense_16 (Dense)	(None, 100)	12250100
dense_17 (Dense)	(None, 100)	10100
dropout_4 (Dropout)	(None, 100)	0
dense_18 (Dense)	(None, 100)	10100
dense_19 (Dense)	(None, 100)	10100
Total params: 12,463,400		
Trainable params: 12,460,600		
Non-trainable params: 2,800		




CNN with Averagepool

Layer (type)	Output Shape	Param #
conv2d_16 (Conv2D)	(None, 297, 297, 100)	2800
average_pooling2d_1 (Average)	(None, 148, 148, 100)	0
conv2d_17 (Conv2D)	(None, 146, 146, 100)	90100
average_pooling2d_2 (Average)	(None, 73, 73, 100)	0
conv2d_18 (Conv2D)	(None, 71, 71, 100)	90100
average_pooling2d_3 (Average)	(None, 35, 35, 100)	0
flatten_6 (Flatten)	(None, 122500)	0
dense_24 (Dense)	(None, 100)	12250100
dense_25 (Dense)	(None, 100)	10100
dropout_6 (Dropout)	(None, 100)	0
dense_26 (Dense)	(None, 100)	10100
dense_27 (Dense)	(None, 100)	10100
Total params: 12,463,400		
Trainable params: 12,460,600		
Non-trainable params: 2,800		




Max VS Average pooling

Max pooling





Average pooling




ResNet50 with large dataset

Accuracy on validation : 16.8%





ResNet50 with
large dataset

Confusion Matrix



ResNet50 with small dataset

Accuracy on validation : 17.8%



Evaluate on test set(ResNet50 on small dataset)

Accuracy top 1: 16.34%

Accuracy majority top 1 : 4.73%

Accuracy top 5 : 33.45%


Accuracy majority top 5 : 15.10%

Accuracy top 10 : 45.40%

Accuracy majority top 10 : 23.73%



ResNet50 with
small dataset

Confusion Matrix



ResNet101

Accuracy on validation : 22.0%



Evaluate on test set(ResNet101)

Accuracy top 1: 17.98%

Accuracy majority top 1 : 4.73%

Accuracy top 5 : 33.68%


Accuracy majority top 5 : 15.10%

Accuracy top 10 : 46.16%

Accuracy majority top 10 : 23.73%



ResNet101

Confusion Matrix



InceptionV3 with sgd

Accuracy on validation : 13.0%



Evaluate on test set(InceptionV3 with sgd)

Accuracy top 1: 10.60%

Accuracy majority top 1 : 4.73%

Accuracy top 5 : 20.74%



Accuracy majority top 5 : 15.10%

Accuracy top 10 : 29.53%

Accuracy majority top 10 : 23.73%

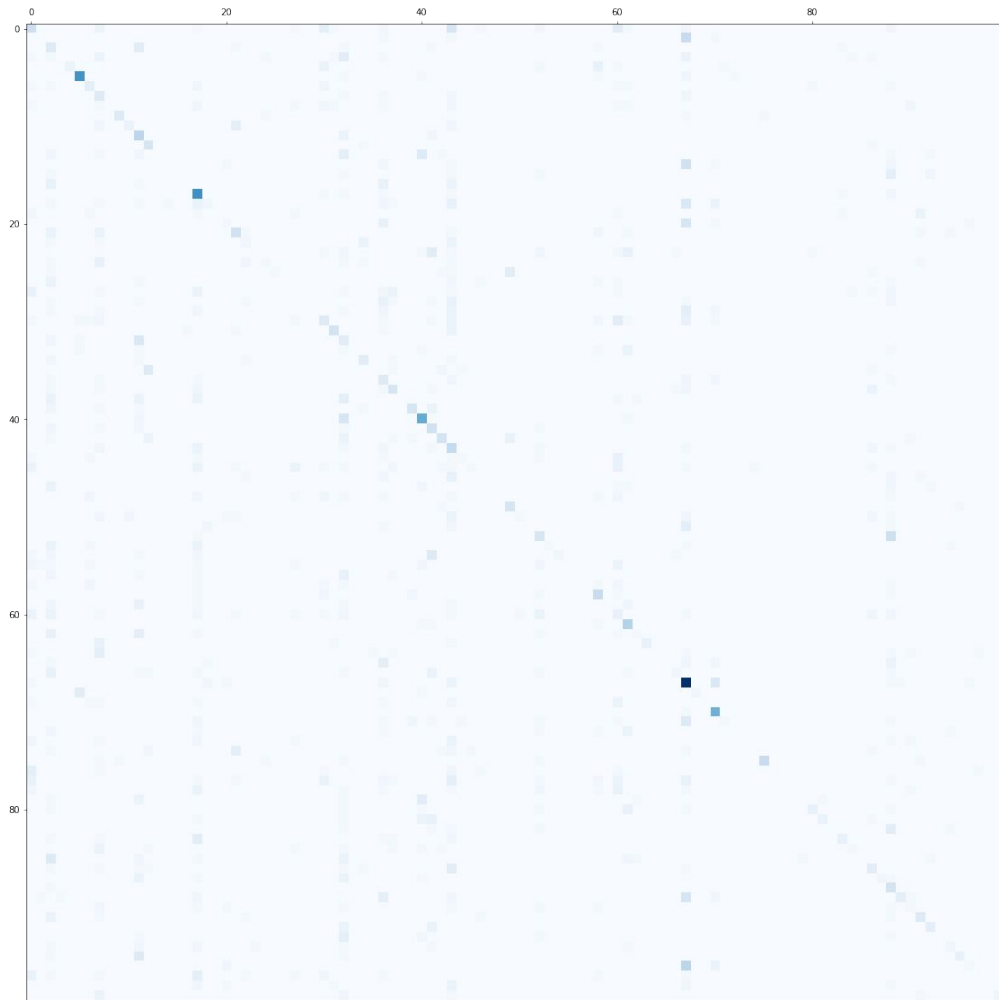
InceptionV3 with rmsprop

Accuracy on validation : 27.8%



InceptionV3
with rmsprop

Confusion Matrix



Evaluate on test set(InceptionV3 with rmsprop)

Accuracy top 1: 29.23%

Accuracy majority top 1 : 4.73%

Accuracy top 5 : 54.48%



Accuracy majority top 5 : 15.10%

Accuracy top 10 : 69.83%

Accuracy majority top 10 : 23.73%


InceptionV3 with adam

Accuracy on validation : 29.9%



InceptionV3
with adam

Confusion Matrix



Evaluate on test set(InceptionV3 with adam)

Accuracy top 1: 28.06%

Accuracy majority top 1 : 4.73%

Accuracy top 5 : 57.23%

Accuracy majority top 5 : 15.10%

Accuracy top 10 : 69.83%



Accuracy majority top 10 : 23.73%



Equal to InceptionV3 with rmsprop

ResNet101 with rmsprop

Accuracy on validation : 37.6%



Evaluate on test set(ResNet101 with rmsprop)

Accuracy top 1: 35.26%

Accuracy majority top 1 : 4.73%

Accuracy top 5 : 64.56%


Accuracy majority top 5 : 15.10%

Accuracy top 10 : 76.92%

Accuracy majority top 10 : 23.73%



ResNet101
with rmsprop

Confusion Matrix





ResNet101 with rmsprop

Loss seems to be decreasing -> increase epoch



ResNet101 with rmsprop (20 epochs)

Accuracy on validation : 40.1%



Evaluate on test set(ResNet101 with rmsprop 20 epochs)

Accuracy top 1: 39.72%

Accuracy majority top 1 : 4.73%

Accuracy top 5 : 67.84%


Accuracy majority top 5 : 15.10%

Accuracy top 10 : 79.44%

Accuracy majority top 10 : 23.73%



ResNet101
with rmsprop
(20 epochs)

Confusion Matrix



ResNet101 with adam(20 epochs)

Accuracy on validation : 38.30%



Evaluate on test set(ResNet101 with adam 20 epochs)

Accuracy top 1: 40.25%

Accuracy majority top 1 : 4.73%

Accuracy top 5 : 68.07%


Accuracy majority top 5 : 15.10%

Accuracy top 10 : 79.44%

Accuracy majority top 10 : 23.73%

ResNet101
with adam
(20 epochs)

Confusion Matrix



Model Summary



Model	Accuracy Top 1	Accuracy Top 5	Accuracy Top 10
ResNet50 on test set	16.34%	33.45%	45.40%
ResNet101 on test set	17.98%	33.68%	46.16%
InceptionV3 w/ SGD	10.60%	20.74%	29.53%
InceptionV3 w/ rmsprop	29.23%	54.48%	69.83%
InceptionV3 w/ adam	28.06%	57.23%	69.83%
ResNet101 w/ rmsprop	35.26%	64.56%	76.92%
ResNet101 w/ rmsprop 20 epochs	39.72%	67.84%	79.44%
ResNet101 w/ adam 20 epochs	40.25%	68.07%	79.44%

Dataset Cleaning



ResNet101 with adam(20 epochs) and cleaned dataset

Accuracy on validation : 35.0%



Evaluate on test set(ResNet101 with adam 20 epochs and cleaned dataset)

Accuracy top 1: 38.25%

Accuracy majority top 1 : 4.73%

Accuracy top 5 : 68.84%


Accuracy majority top 5 : 15.10%

Accuracy top 10 : 79.50%

Accuracy majority top 10 : 23.73%

ResNet101
with adam and
cleaned dataset
(20 epochs)

Confusion Matrix



Summary

- Adam and rmsprop is more preferable than sgd in this case.






Image Credit: <https://cs231n.github.io/neural-networks-3/>

ResNet101 with sgd(30 epochs)

Accuracy on validation : 29.6%



Evaluate on test set(ResNet101 with sgd 30 epochs and cleaned dataset)

Accuracy top 1: 27.30%

Accuracy majority top 1 : 4.73%

Accuracy top 5 : 54.19%


Accuracy majority top 5 : 15.10%

Accuracy top 10 : 66.20%

Accuracy majority top 10 : 23.73%

ResNet101
with sgd
(30 epochs)

Confusion Matrix



ResNet101(sgd) 10 vs 30 epochs

10 epochs

Accuracy top 1: 17.98%
Accuracy majority top 1 : 4.73%

Accuracy top 5 : 33.68%
Accuracy majority top 5 : 15.10%

Accuracy top 10 : 46.16%
Accuracy majority top 10 : 23.73%

30 epochs

Accuracy top 1: 27.30%
Accuracy majority top 1 : 4.73%

Accuracy top 5 : 54.19%
Accuracy majority top 5 : 15.10%

Accuracy top 10 : 66.20%
Accuracy majority top 10 : 23.73%



Let's try 100 epochs(with callback)

But...

our callback function stop at the 60th epoch.

ResNet101 with sgd(60 epochs)

Accuracy on validation : 35.1%



Evaluate on test set(ResNet101 with sgd 60 epochs and cleaned dataset)

Accuracy top 1: 33.86%

Accuracy majority top 1 : 4.73%

Accuracy top 5 : 60.98%


Accuracy majority top 5 : 15.10%

Accuracy top 10 : 74.17%

Accuracy majority top 10 : 23.73%

ResNet101
with sgd
(60 epochs)

Confusion Matrix



ResNet101(sgd) 30 vs 60 epochs

30 epochs

Accuracy top 1: 27.30%
Accuracy majority top 1 : 4.73%

Accuracy top 5 : 54.19%
Accuracy majority top 5 : 15.10%

Accuracy top 10 : 66.20%
Accuracy majority top 10 : 23.73%

60 epochs

Accuracy top 1: 33.86%
Accuracy majority top 1 : 4.73%

Accuracy top 5 : 60.98%
Accuracy majority top 5 : 15.10%

Accuracy top 10 : 74.17%
Accuracy majority top 10 : 23.73%

sgd vs adam vs rmsprop on ResNet101

optimizer/epochs	Accuracy top1	Accuracy top5	Accuracy top10
rmsprop/10	35.26	64.56	76.92
rmsprop/20	39.72	67.84	79.44
adam/20	40.25	68.07	79.44
adam/20 cleaned	38.25	68.84	79.50
sgd/30 cleaned	27.30	54.19	66.20
sgd/60 cleaned	33.86	60.98	74.17

Garbage in Garbage out

- The machine is as intelligent as the data/features we put in
- “Garbage in, Garbage out”
- Data cleaning is often done to reduce unwanted things

Ref: https://github.com/ekapolc/pattern_2019

